

Announcements

- PP2 is released, due March 5th
- Midterm 1 is on March 3rd

Topic 3: Graphs

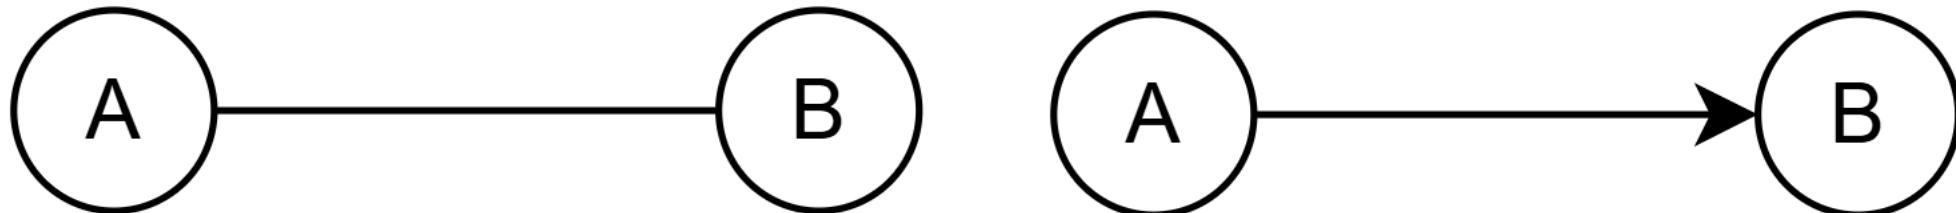
By Professor Hudson Lynam

Graph Definition

- Definition: *Graph*
- *A graph $G=(V,E)$ where V is a finite non-empty set of vertices (a.k.a. nodes) and E is a binary relation on V .*
- Note that definitions of “graph” vary: this is a common definition we’ll be using for this class

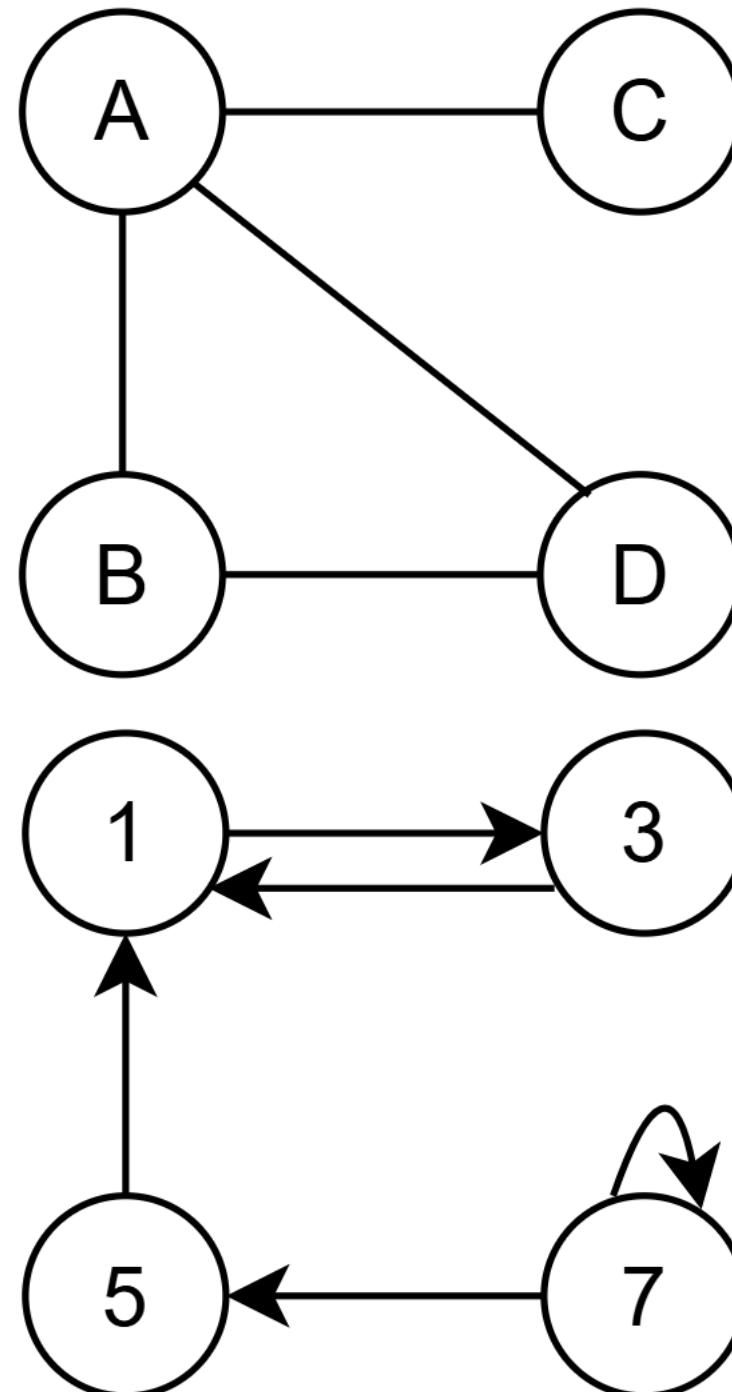
Undirected versus Directed Graphs

- Undirected:
 - The pairs of vertices in E are unordered. In other words, each edge is a set, like $\{A, B\}$.
 - When we say “ G is a graph,” the default assumption is that G is undirected
- Directed
 - Each edge has an order, like a tuple (A, B) .
 - Because directed graphs are common, we usually abbreviate to “digraph.”



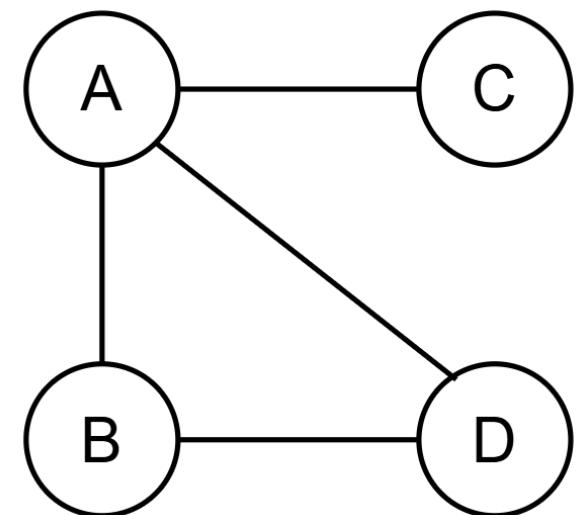
Graphs Example

- What would V be for this graph?
- $V=\{A, B, C, D\}$
- $E=\{\{A,C\}, \{A,B\}, \{A,D\}, \{B, D\}\}$
- What about V and E for this digraph?
- $V=\{1, 3, 5, 7\}$
- $E=\{(1,3), (3,1), (5,1), (7, 7), (7, 5)\}$
- Note: Don't use bidirectional arrows for undirected graphs, and always use arrows to show two edges between vertices in a digraph



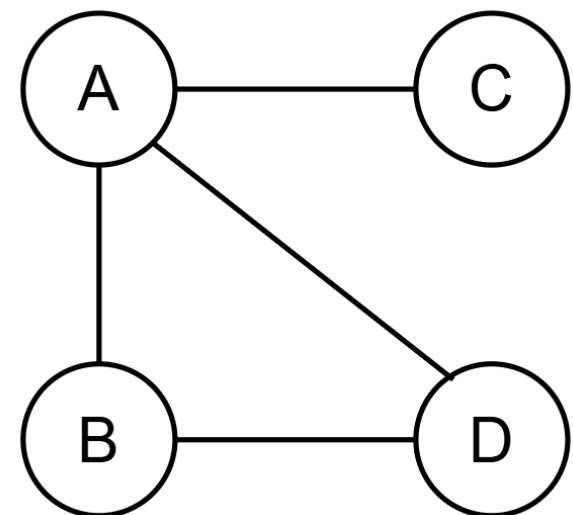
Graph Terminology

- Definition: *Adjacent Vertices*
- *2 vertices are adjacent when an edge exists between them*
- In our example graph, C and D are not adjacent, B and C are not adjacent, every other pair of vertices are adjacent
- Definition: *Path (a.k.a. Walk)*
- *A path is a sequence of adjacent vertices*
- Example: C-A-B-A-D-B.
- A simple path's vertices appear exactly once. The above example is *not* simple. What would be a simple path here?



Graph Terminology

- Definition: *Cycle*
- *A cycle is a path of length ≥ 1 that begins and ends with the same vertex*
- In our example graph, A-C-A is a cycle. D-B-A-C-A-D is a cycle.
- A cycle is **simple** if all the vertices in the path only occur once (except for the first and last vertices)
- Notes: “the cycle ABD” means A-B-D-A
- Common assumption: simple cycles in undirected graphs must have length ≥ 3




Graph Terminology

- Definition: *Incident Edges*
- *In any graph, an edge is incident on both participating vertices*
- *In digraphs, an edge is **incident from** the source and **incident to** the destination vertex*

- Definition: *Degree*
- *The degree of a vertex is its # of incident edges*

- Note: a self-loop adds to a vertex's degree
- What is the degree of 1? In versus out?
- In-degree of 1: 2, out-degree of 1: 1

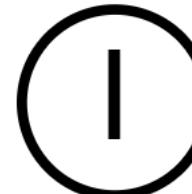
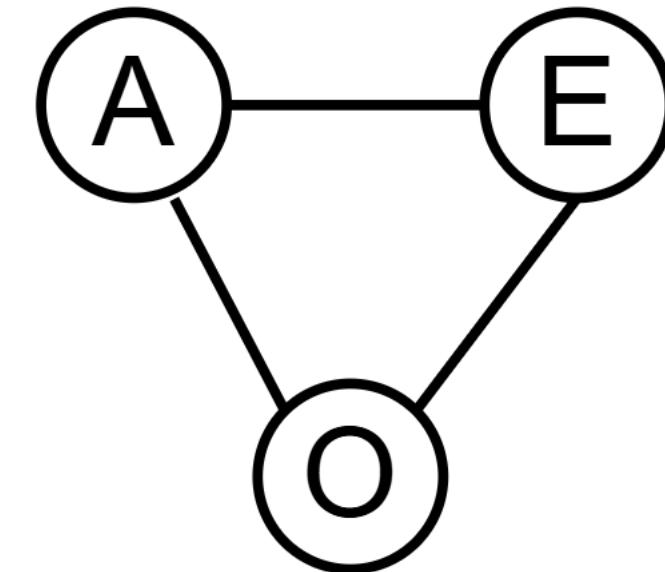


Graph Representations

- How can we store all of a graph's vertices and edges? Three common representations:
 1. Edge Lists
 2. Adjacency Matrices
 3. Adjacency Lists
- Which should we choose? Depends on the operations we want to be efficient.

Edge Lists

- An edge list is a collection of edges within a graph. Example:
- I is an isolate in this graph. Obviously, not an ideal representation for graphs with isolates
- Note: In a digraph, just allow 1st vertex in each edge to be the source or “from” vertex

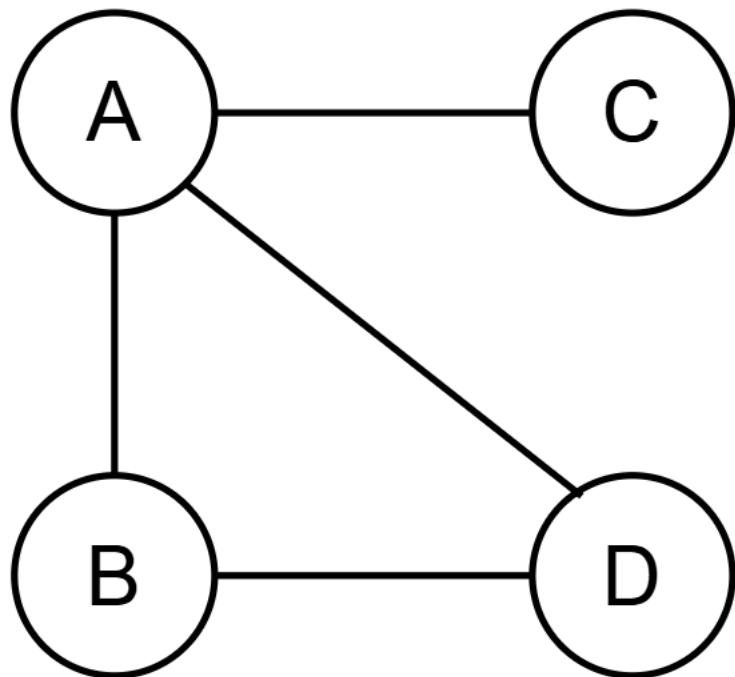


Adjacency Matrices

- Adjacency Matrices are a good choice when the quantity of vertices is bounded and the graph has a relatively large quantity of edges.
- The **adjacency matrix** of a graph $G=(V, E)$ is a $|V| \times |V|$ matrix M in which element $M[i][j] = 1$ iff $\{v_i, v_j\}$ or $(v_i, v_j) \in E$. Otherwise, $M[i][j] = 0$.

Adjacency Matrices

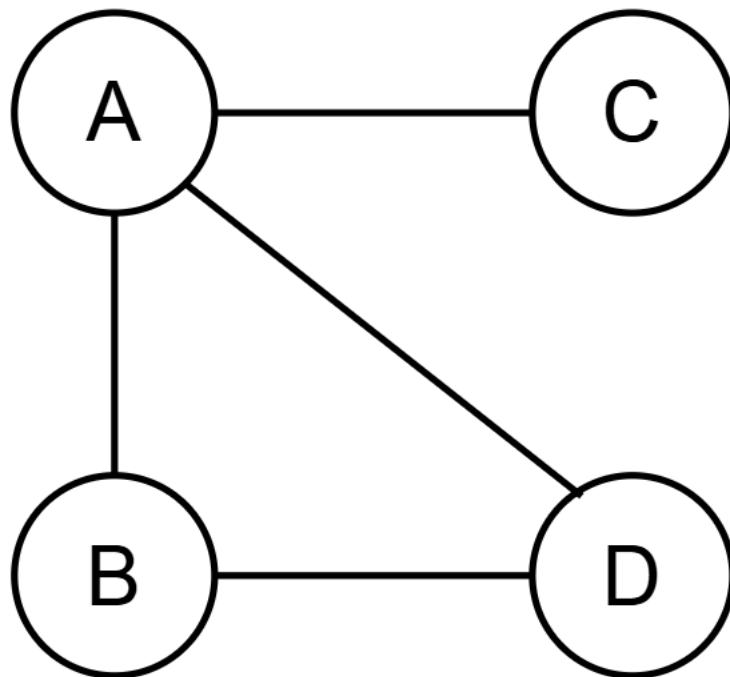
- Let's take a look at our example graph:



	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	1	1	0	0

Adjacency Matrices

- Did we even need the full matrix? Could we represent the edges more efficiently? How?

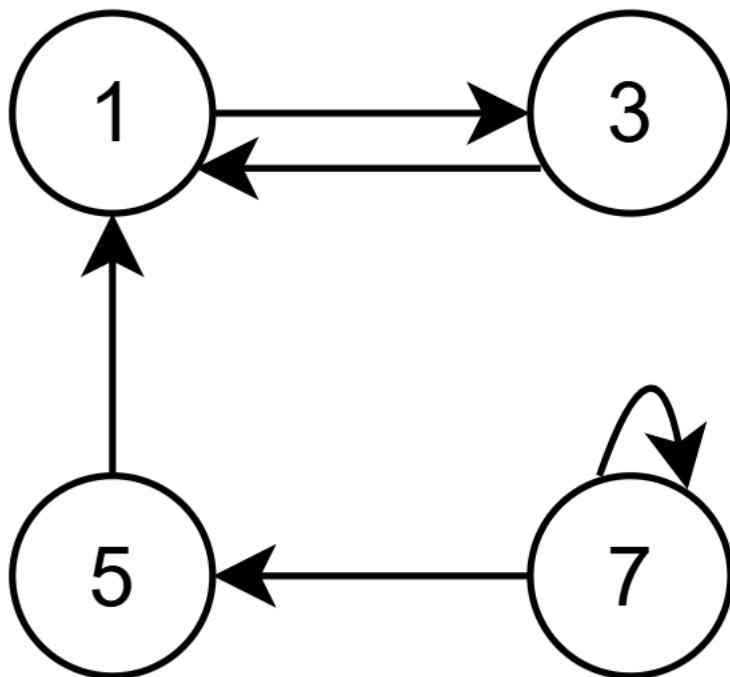


	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	1	1	0	0

- Note: this is a lower left triangular matrix

Adjacency Matrices

- Try to make an adjacency matrix for this digraph. Cols are To, Rows are From



	1	3	5	7
1	0	1	0	0
3	1	0	0	0
5	1	0	0	0
7	0	0	1	1

Adjacency Matrices

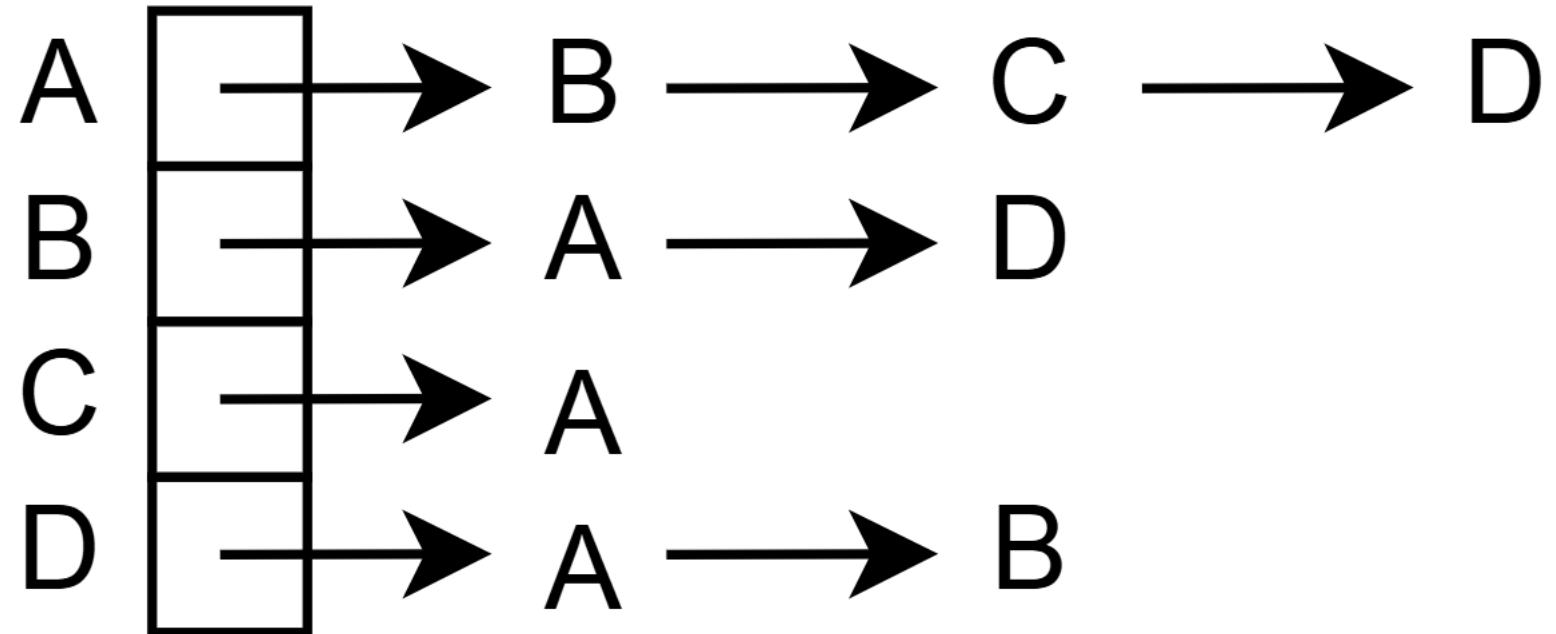
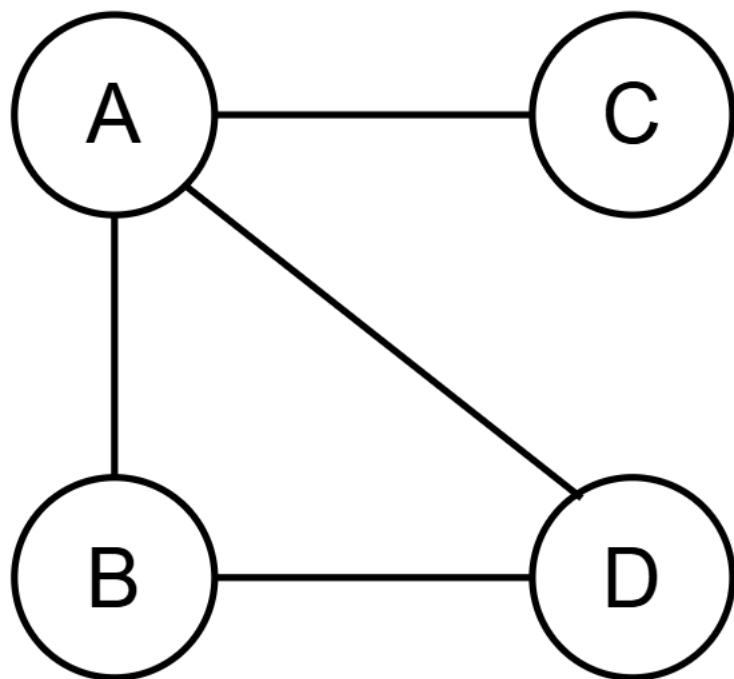
- Adjacency matrices make many operations easy. For example...
- How would we find the degree of a vertex in an AM?
- Count the row or col for the vertex if it's a graph, count the row for out-degree in a digraph, or col for the in-degree
- Which edges are incident to a given vertex?
- Same solution as finding degree!
- But there's a downside:
- We need $O(|V|^2)$ space to store an AM. Which is why AM is a good choice when the number of vertices is low

Adjacency Lists

- The **adjacency list** of a graph $G=(V, E)$ is an array of $|V|$ lists of vertices, with each list holding the vertices to which the vertex is adjacent (for a graph) or the vertices incident from a vertex (for a digraph).

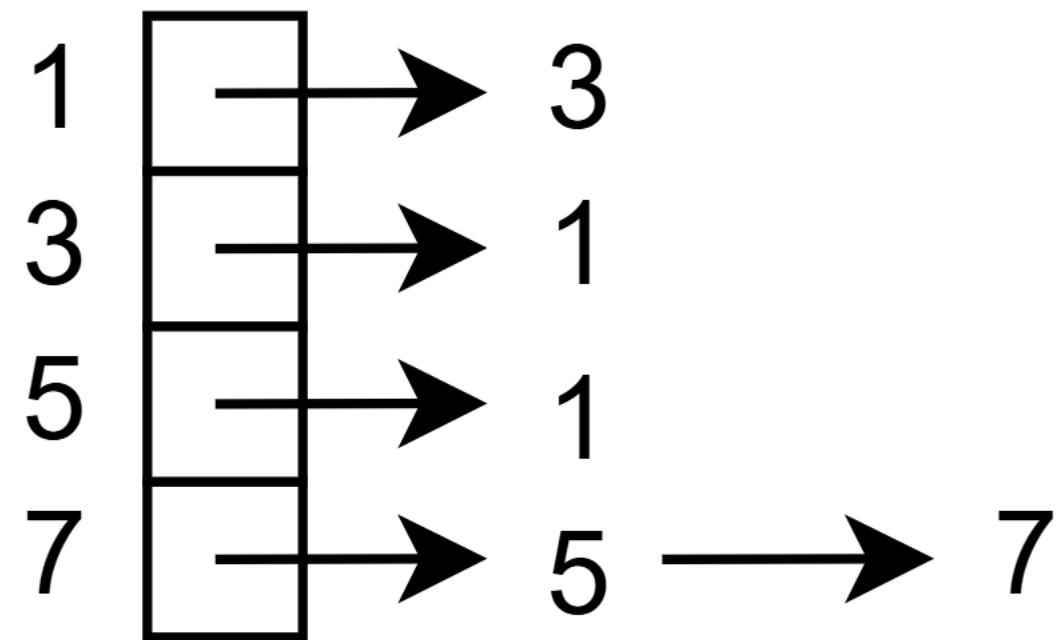
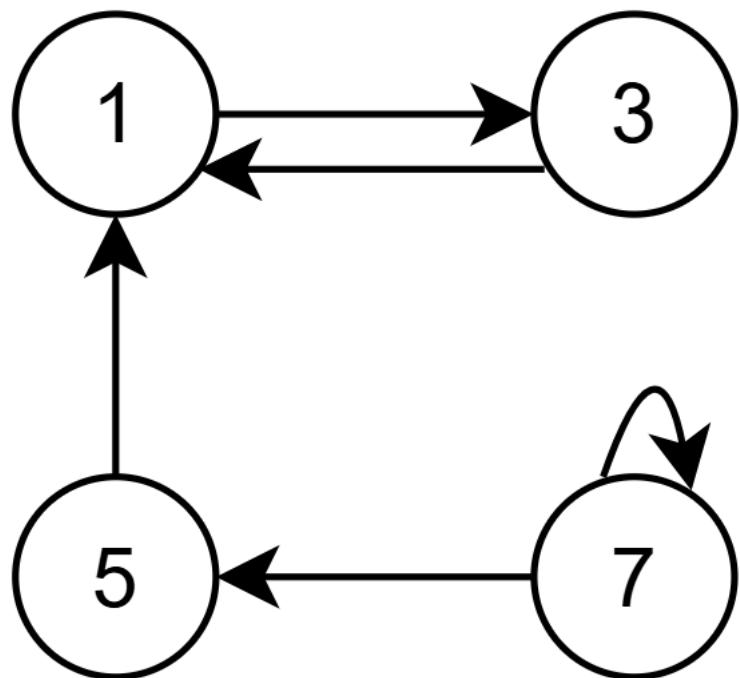
Adjacency Lists

- Let's take a look at our example graph:



Adjacency Lists

- Now go ahead and try to construct the AL for our digraph example.



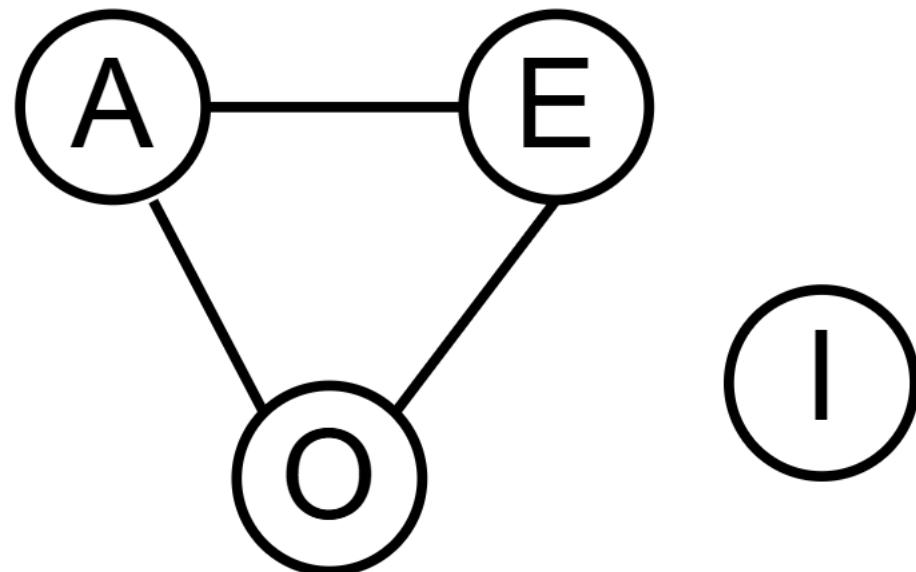
Adjacency Lists

- ALs are efficient for some operations, but not others
- How about finding the degree for a given vertex?
- Easy for graphs, but digraphs require searching the entire AL!
- What's the storage requirement for AL?
- $O(|V|+|E|)$; this is an example of having 2 instance characteristics!
- In summary, when choosing a graph representation, consider the graph (# of vertices vs edges) and the operations you expect to do on the graph (and how expensive they're likely to be)

Connectedness of Graphs

- Definition: *Connected Graph*
- *A graph G is connected when every pair of vertices in G is connected.*


- Is this graph connected?
- How about this graph?
- You can't get to I from any other vertex, so no, not a connected graph



Connectedness of Digraphs

- Definition: Weakly Connected Digraph
- A digraph G is a weakly connected digraph if it is connected when edge direction is ignored.

- Definition: Strongly Connected Digraph
- A digraph G is a strongly connected digraph if, for every pair of vertices there is a path between them in both directions.



Determining that a Graph is Connected

- What are some of the tree traversals you've learned?
- Preorder, Postorder, Inorder, and Level Order
- Would any of these work for graphs?
- No; because graphs can have cycles
- We need a different way to traverse (or search through) graphs!

Breadth-First Search (BFS)



Breadth-First Search (BFS)

- Like the spread of liquid after popping a water balloon, BFS starts at some “center” and fans in all directions
- BFS relies on a queue

BFS:

Enqueue the source vertex

While the queue is not empty:

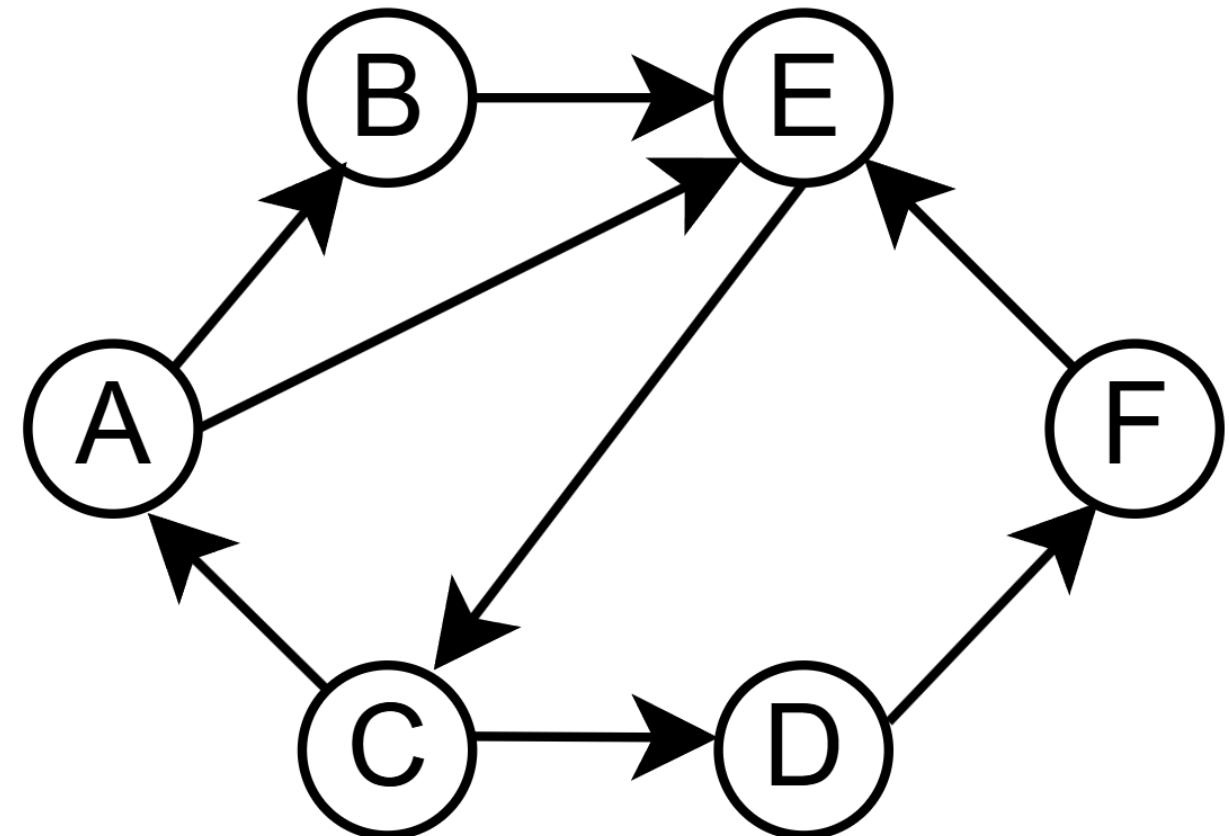
 Dequeue the first vertex

 Act on the vertex

 Enqueue the adjacent but previously undiscovered vertices

Breadth-First Search (BFS)

- Let's do a BFS on this graph, with A as the source vertex
- Queue: A, prev seen: A. Dequeue A.
- Queue: B, E prev seen: A, B, E. Dequeue B.
- Queue: E, prev seen: A, B, E. Dequeue E.
- Queue: C, prev seen A, B, E, C. Dequeue C.
- Queue D, prev seen all but F. Dequeue D.
- Queue: F, prev seen all. Dequeue F, queue is empty, finished.



Breadth-First Search (BFS)

- Some notes on BFS:
- If vertices and edges are not reachable from the source vertex, BFS cannot find them.
 - To deal with this, we can re-apply BFS to the remaining portions of the graph to find additional connected components of the graph
- How efficient is BFS? Depends on the graph representation.
Discuss: Adjacency Matrix versus Adjacency List?
- AM: $O(|V|^2)$
- AL: $O(|V| + |E|)$

Depth-First Search (DFS)

- The idea is: when we see a new vertex, we visit that vertex!
- DFS relies on a stack

DFS:

Push the source vertex

While the stack is not empty:


 Pop the first vertex

 Act on the popped vertex

 Push its adjacent but previously undiscovered vertices

Depth-First Search (DFS)

- Let's do a DFS on this graph, with A as the source vertex
- Push: A, prev seen: A. Pop A.
- Push: B, E prev seen: A, B, E. Pop E.
- Push: C, prev seen: A, B, E, C. Pop C.
- Push: D, prev seen A, B, E, C, D. Pop D.
- Push: F, prev seen all. Pop F.
- Pop: B. Stack is empty, finished.



Depth-First Search (DFS)

- Some notes on DFS, first like BFS:
- If vertices and edges are not reachable from the source vertex, DFS cannot find them.
 - To deal with this, we can re-apply DFS to the remaining portions of the graph to find additional connected components of the graph
- Unlike BFS: DFS is easy to code recursively.
- How efficient is DFS? Depends on the graph representation. Discuss: Adjacency Matrix versus Adjacency List?
- AM: $O(|V|^2)$
- AL: $O(|V| + |E|)$

More Graph Terminology

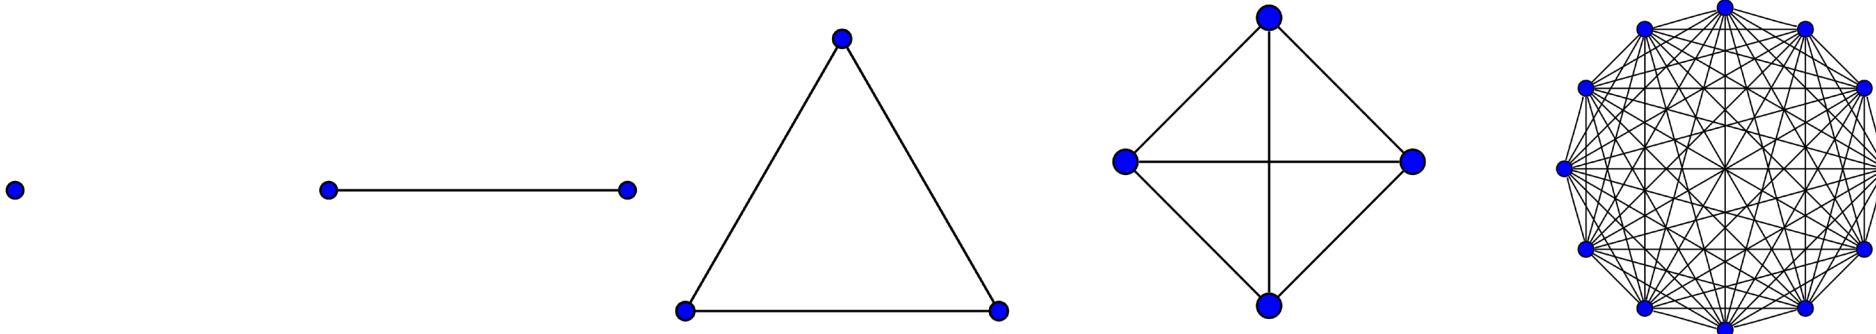
- Definition: Simple *Graph*
- *A simple graph is a graph w/o self-loops and w/o multiple edges between a pair of vertices*

- Definition: Acyclic *Graph*
- *An acyclic graph is a graph with no cycles.*
- *In undirected acyclic graphs, we'll ignore non-simple cycles*
- *A directed acyclic graph is often called a “DAG”*

More Graph Terminology


- Definition: Subgraph
- G' is a subgraph of G if $V' \leq V$, $E' \leq E$, and E' uses only the vertices in V'

- Definition: Complete Graph
- A complete graph is an undirected graph in which all pairs of n vertices are adjacent is a complete graph and is denoted K_n



More Graph Terminology

- Definition: Bipartite *Graphs*
- *A graph G whose vertices are partitionable into two subsets V_1 and V_2 such that all of G's edges connect a vertex from V_1 to one from V_2 is a bipartite graph.*



More Graph Terminology

- Definition: Weighted *Graph*
- *Any graph whose edges have assigned costs weighted graph. (On an unweighted graph, edges have the same weight (usually 1).)*

- Definition: Forest
- *A forest is a simple undirected acyclic graph*

- Definition: Tree
- *A tree is a connected simple undirected acyclic graph*

The Single-Source Shortest Path Problem


- After class, say you want to walk to Gould-Simpson taking the fewest # of steps. What route do you take?
- This is an example of the Single-Pair Shortest Path (SPSP) problem.
- A more general problem is the Single-Source Shortest Path problem: what is the shortest path from the start to all destinations?
- We would like to solve this problem for...
 - Weighted and unweighted graphs
 - Directed and undirected graphs

Negative Edge Weights

- In some cases, negative edge weights make sense:
- Example 1: Elevation changes. A graph of cities relative elevation levels would have negative edge weights for edges from higher elevation cities (Denver) to lower elevation cities (Tucson)
- Example 2: Costs versus Profit. In a graph of items you want to sell versus items you want to buy. Positive edge weights might represent money gained, negative edge weights might represent money lost


Negative Edge Weights

- Can a shortest path include a cycle?
- If the cycle weight is >0 , the cycle is skipped
- If the cycle weight is $=0$, cycle *may* be included
- If the cycle weight is <0 , we'll *never* leave it!



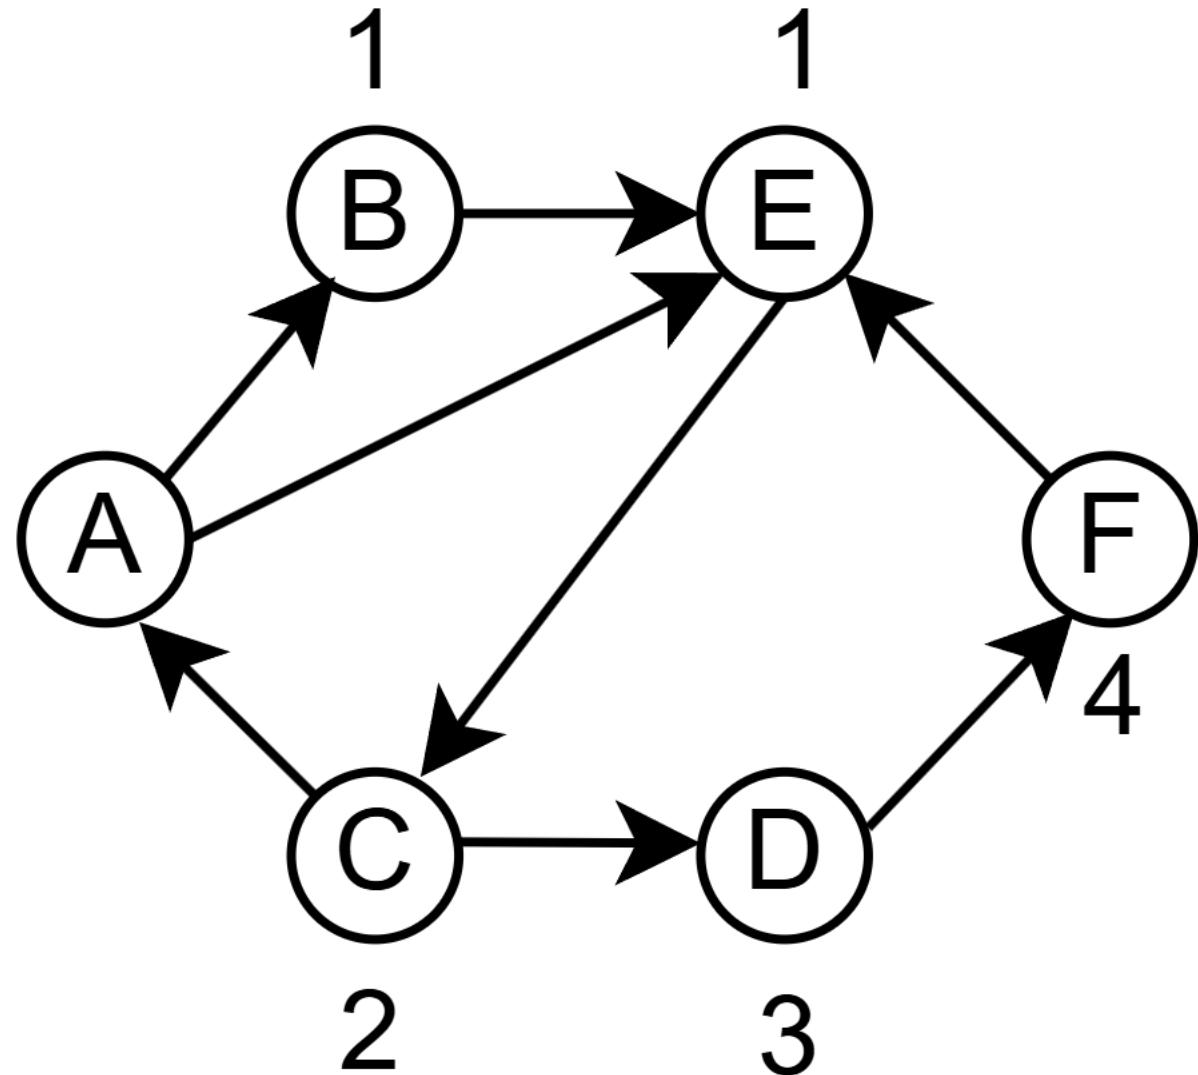
SSSP

- In a simple unweighted graph, finding the shortest paths from the source is easy! What could we use to do it, and how?
- BFS! Update when we reach each vertex



SSSP

- A path with more edges can be cheaper, in a **weighted** graph
- What if the edge weight from A to B was 1, and B to E also 1, but A to E was weight 5?
- BFS won't work, because we need to update the costs of paths as we search
- We need a different algorithm...



Dijkstra's (SSSP) Algorithm

- Dijkstra's Algorithm relies on four data structures:
- $w(x, y)$: The weight of the edge connecting vertices x and y .
- $d(a, b)$: The path cost (distance) from vertex a to vertex b .
- Fringe: The set of vertices that we know we can reach from the source
- Known: The set of vertices whose shortest paths from the source are known.

Dijkstra's (SSSP) Algorithm

Dijkstra's:

$d(\text{source}, \text{source}) = 0$

$d(\text{source}, x) = +\infty$, for all non-source vertices

Known \leftarrow source

Fringe \leftarrow vertices adjacent to source

While $|\text{Fringe}| \neq 0$:

f \leftarrow fringe vertex with smallest $d(\text{source}, f)$

Move f from Fringe to Known


Add unKnown, unFringe vertices that are adjacent to f to the Fringe

Update Fringe content as necessary

- Notes:
- $d(\text{source}, f) = d(\text{source}, t) + w(t, f)$ where $t \in \text{Known}$
- Must update Fringe when another path to a vertex is found.

Dijkstra's (SSSP) Algorithm

- Let's try it with Source=D.
Notation $V_{\#,P}$ where V is a vertex, # = $d(\text{source}, v)$, and p = immediate previous vertex



Dijkstra's (SSSP) Algorithm

Known

$D_{0,-}$

$H_{2,D}$

$E_{3,D}$

$C_{4,H}$

$G_{6,H}$

$B_{7,C}$

$F_{8,G}$

$A_{9,F}$

Fringe

$H_{2,D} E_{3,D}$

$E_{3,D} C_{4,H} G_{6,H}$


$C_{4,H} G_{6,H}$

$G_{6,H} B_{7,C}$

$B_{7,C} F_{8,G}$

$F_{8,G} A_{10,B}$

$A_{9,F}$



Dijkstra's (SSSP) Algorithm

- How efficient is this algorithm? Both in terms of space, and execution?
- Space: $O(|V|)$ (in addition to the graph representation, AM or AL)
- Execution: Complicated!
- $O(|V|^2 + |E|)$ when Fringe is a list
- $O(|V| \log_2 |V| + |E|)$ if Fringe is represented with a Fibonacci Heap
- Final Note: will this algorithm work with negative edge weights?
- No! Once a vertex has been marked as Known, its current distance from source is set. If a negative weight edge is found that would reduce the distance to a Known vertex, the algorithm can't use it


Dijkstra's Ex. 2

Known

$F_{0,-}$
 $E_{4,F}$
 $D_{6,E}$
 $G_{6,E}$
 $C_{8,D}$
 $A_{9,D}$
 $B_{14,A}$

Fringe

$E_{4,F}$
 $D_{6,E} G_{6,E}$
 $G_{6,E} C_{8,D} A_{9,D}$
 $C_{8,D} A_{9,D}$
 $A_{9,D} B_{15,C}$
 $B_{14,A}$



Bellman-Ford (SSSP) Algorithm

- Like Dijkstra's: Bellman-Ford is an SSSP algorithm
- Unlike Dijkstra's: Bellman-Ford can handle negative edge weights!
- Notes about algorithm:
 - $w(x, y)$ and $d(a, b)$ are the same as in Dijsktra's
 - There is a loop that continues $|V| - 1$ times... then a final loop, the “bonus round” that checks for negative edge-weight cycle

Bellman-Ford (SSSP) Algorithm

Bellman-Ford:

$d(\text{source}, \text{source}) = 0$

$d(\text{source}, x) = +\infty$, for all non-source vertices

Loop $|V| - 1$ times:

 For each edge (x, y) in E :

 if $d(\text{source}, x) + w(x, y) < d(\text{source}, y)$

$d(\text{source}, y) \leftarrow d(\text{source}, x) + w(x, y)$

For each edge (x, y) in E :

 if $d(\text{source}, x) + w(x, y) < d(\text{source}, y)$

 Report failure due to a negative edge-weight

Report successful completion

Bellman-Ford (SSSP) Algorithm

$d(\text{source}, *)$ From

$P \quad 0$ -

$Q \quad inf$ -

$R \quad inf$ -

$d(\text{source}, *)$ From

$P \quad 0$ -

$Q \quad 1$ P

$R \quad inf$ -

$d(\text{source}, *)$ From


$P \quad 0$ -

$Q \quad 1$ P

$R \quad 3$ Q

Edges	Pass 1	Pass 2	Bonus
$R \rightarrow Q$	---	---	$d=0$
$R \rightarrow P$	---	---	<i>Report failure</i>
$Q \rightarrow R$	---	$d=3$	
$P \rightarrow Q$	$d=1$	$d=1$	

if $d(\text{source}, x) + w(x, y) < d(\text{source}, y)$
 $d(\text{source}, y) \leftarrow d(\text{source}, x) + w(x, y)$



Bellman-Ford (SSSP) Algorithm

- Questions about Bellman-Ford:
- Why does the first loop iterate $|V| - 1$ times?
 - To ensure that we reach all reachable vertices!
- Can we consider the edges in any order?
 - Yes! Order does not matter for Bellman-Ford
- Could a negative edge-weight cycle be found sooner?
 - Not without augmenting the base algorithm


Bellman-Ford (SSSP) Algorithm

- Bellman-Ford efficiency analysis, or... What does this negative edge-weight handling capability cost us?
- Initialization: $O(|V|)$ (The $d()$ matrix initialization)
- Nested Loops:
 - Outer loop: $O(|V|)$ (“Loop $|V| - 1$ times”)
 - Inner loop: $O(|E|)$ (“For each edge...”)
 - Combined: $O(|V|^*|E|)$
- Bonus Round: $O(|E|)$ (one more run of inner loop)
- Total: $O(|V| + |V|^*|E| + |E|)$, or...
- $O(|V|^*|E|)$

What about all pairs shortest path?

- We have Dijkstra's and Bellman-Ford for SSSP, but what about the shortest distances for all pairs of vertices? What would a naïve solution be?
- Use an SSSP algorithm on every single vertex... But surely we can do better?
- Yes, we can: Floyd-Warshall's algorithm... Which we'll cover in a later topic!

Topological Sorting



- Consider this piece of the CS course prerequisite graph. What would a legal one-course-per semester class schedule be? In other words, what would a linear ordering of the vertices such that the edge restrictions are respected be?
- Well, we can look at it, and eyeball a solution... but do we already have an algorithm that might be able to do this?

Topological Sorting

DFS (recursive variant):

Given a source vertex

Loop:


 Select an unvisited neighbor of the source

 Call DFS(neighbor)

Return


- How can modify this DFS variant to topologically sort a graph?
- We could prepend the source to the ordered list of vertices... After the recursive call!
- We would need to: create an initial empty list, mark vertices as having been visited, and restart the sort on any unreachable vertex

Topological Sorting



- Source: 210, pick a neighbor: 345
- Pick a neighbor? Can't. Prepend 345 to the list and backtrack.
- Pick: 335. Pick a neighbor? Can't. Prepend 335 to the list.
- Back at 210, pick a neighbor? Can't. Prepend 210 to the list.
- List so far: 210 335 345

Topological Sorting



- Source: 120, pick a neighbor: 244
- Pick a neighbor? Can't. Prepend 244 to the list and backtrack.
- Back at 120, pick a neighbor? Can't. Prepend 120 to the list.
- Source: 144, pick a neighbor? Can't. Prepend 144 to the list.
- Source: 110, pick a neighbor? Can't. Prepend 110 to the list.
- Final List: 110 144 120 244 210 335 345


Topological Sorting

- What is the big-O for DFS?
- For AM: $O(|V|^2)$
- For AL: $O(|V| + |E|)$
- What cost(s) are added by Topological Sorting?
- Prepending to a Linked List is $O(1)$, so $O(|V|)$
- Asymptotically, is there a change in the Big-O from DFS?
- Nope! Same classification, even though there's a bit more work

Spanning Trees

- Definition: *Spanning Tree*
- *A spanning tree of a connected graph G is a connected subgraph of G that is a tree and includes all of G's vertices. (The tree is thus forced to include only |V| - 1 of G's edges.)*

- What would a possible spanning tree of this graph be?
- Remove any single one of the three edges, and the graph becomes a spanning tree



Spanning Trees

- Why $|V|-1$ edges?
- It's the minimum number of edges needed to tie all of the vertices together. Any more creates a cycle
- Can we find spanning trees in digraphs?
- Possibly, depending on the available edges. Typically, we do spanning trees on undirected graphs.
- Can we have a ‘spanning forest’?
- Yes! We don’t worry about them for this course, but it is another type of spanning tree application
- Cayley’s formula: A complete graph has $|V|^{|V|-2}$ spanning trees

Minimal Cost Spanning Trees

- Definition: Minimal Cost *Spanning Tree (MCST)*
 - *A minimal cost spanning tree of a connected graph G is one of the spanning trees of G that has the minimum total edge weight of all spanning trees of G .*
-
- In an unweighted graph, any spanning tree is the graph's MCST
 - We'll talk about two classic algorithms that find MCSTs: Prim's and Kruskal's algorithms

Prim's Algorithm

- The idea is to grow a single tree, one edge at a time
- The algorithm:

Prim's:

Initialize the MCST, call it T , with one $v \in V$

Repeat:


 Find a lowest-cost edge such that $T \cup e$ is a tree

$T \leftarrow T \cup e$

Until T contains $|V| - 1$ edges


Prim's Algorithm

- Let's start at B. What would the first edge in our MCST be?
- AB! Next?
- AC
- CD or BD... Let's go with CD
- DE
- CF
- Could've gone with BD—that would've given us the other possible MCST for this graph



Prim's Algorithm

- Let's start at B. What would the first edge in our MCST be?
- AB! Next?
- AC
- CD or BD... Let's go with CD
- DE
- CF
- Could've gone with BD—that would've given us the other possible MCST for this graph



Kruskal's Algorithm

- The idea is to build a forest and connect it into a single tree
- The algorithm:

Kruskal's:

Create a forest F of all vertices in G

Repeat:

 Select a lowest-cost edge e from E

 If adding e to F would create a cycle:

 remove e from consideration


 Else:

 Add e to F

Until $|V| - 1$ edges have been added


Kruskal's Algorithm

- Same graph, initialize F with V . What edge would be added first?
- AC
- DE
- AB
- BD (could've been CD)
- CF



Kruskal's Algorithm

- Same graph, initialize F with V . What edge would be added first?
- AC
- DE
- AB
- BD (could've been CD)
- CF



Efficiencies of Prim's and Kruskal's

- Prim's, using a basic implementation:
- $O(|V|^2)$ (Using an adjacency matrix and linear search)
- Can be improved to: $O(|E| + |V| \log_2 |V|)$ using an adjacency list and a Fibonacci Heap
- Kruskal's, using a basic implementation:
- $O(|V|^2)$ (Using an adjacency matrix and linear search)
- Can be improved to: $O(|E| \log_2 |E|)$ using a complex sorting algorithm